



Together for Tomorrow!  
**Enabling People**

Education for Future Generations

# Samsung Innovation Campus

## Artificial Intelligence Course

Chapter 9.

# Deep Learning - Part II

AI Course

Samsung Internal Use Only

# Update History

| Version | Update Details   |
|---------|--|
| v1.0    | <ul style="list-style-type: none"><li>1<sup>st</sup> draft completed</li></ul>   |
| v1.01   | <ul style="list-style-type: none"><li>'Guide for Instructors' page inserted</li></ul>  |
| v1.02   | <ul style="list-style-type: none"><li>Front and Cover page layout updated</li><li>'Coding Exercise' replaced 'Coding Practice'</li><li>'Quiz.' replaced 'Coding Problem'</li></ul> |
| v1.03   | <ul style="list-style-type: none"><li>Samsung, SIC logo updated</li></ul>  |
| v1.04   | <ul style="list-style-type: none"><li>Chapter info. updated</li></ul>  |
|         |  |
|         |  |
|         |  |
|         |  |

## Deep Learning - Part II

UNIT 3.-----

# Deep Learning with Keras

# UNIT 3.

## Deep Learning with Keras

### What this unit is about:

- Building artificial intelligence models with the Keras library.
- Utilization of the DNN, CNN, RNN and LSTM networks.
- Embedding representation of words.
- Natural language processing with the Keras library.

### Expected outcome:

- Proficiency in building artificial intelligence models.

### How to check your progress:

- Coding Exercises.
- Quiz.

# Deep Learning - Part II

## UNIT 3. Deep Learning with Keras.

### 3.1. Keras Basics.

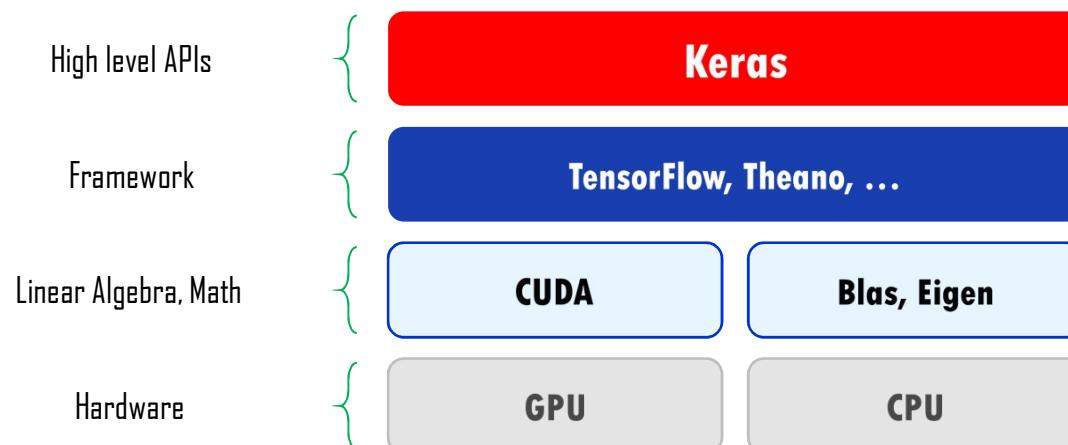
3.2. AI with Keras.

3.3. Natural Language Processing with Keras.

# Keras Basics (1/9)

## About the Keras library:

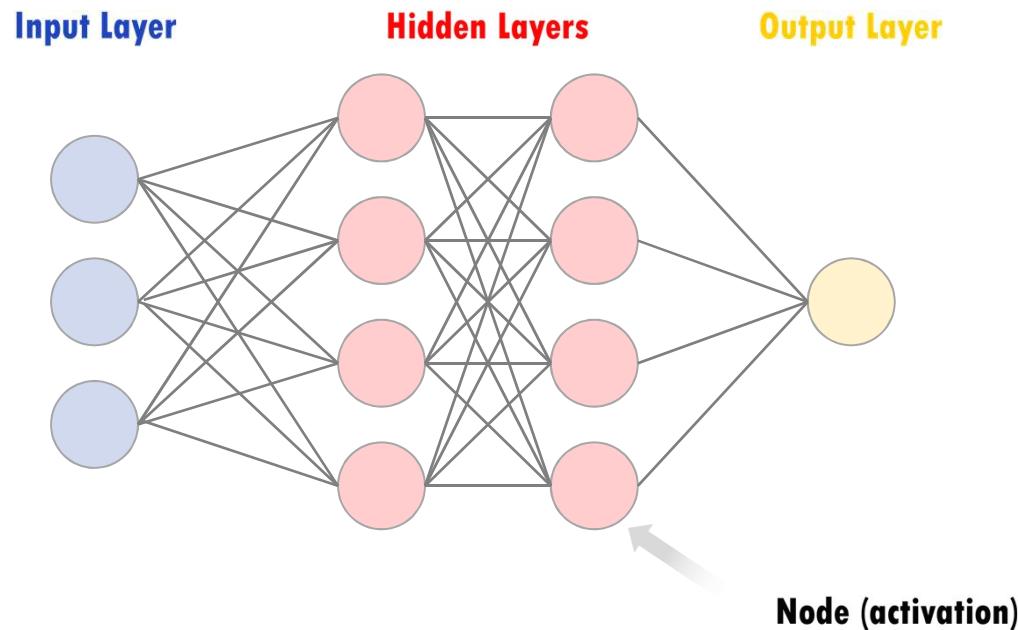
- A high-end deep learning library.
- Does not support the low-end tensor operations.
- Runs on top of Tensorflow.



# Keras Basics (2/9)

## | About the Keras library:

- ▶ It is easy to model deep neural networks with Keras.



# Keras Basics (3/9)

| Training work flow:

Define a model's layer structure (specify the activation functions).

Define the loss function and the optimization method.

Compile the model.

Specify the batch learning method and start the training (gradient descent).

# Keras Basics (4/9)

I Loss (cost) functions in Keras:

| Name                         | Formula   | Keras Expression                               |
|------------------------------|---|--|
| Squared Error<br>(L2 Error)  | $L = \sum_{i=1}^n (y_i - \hat{y}_i)^2$  | <code>loss = "mse"</code>                      |
| Absolute Error<br>(L1 Error) | $L = \sum_{i=1}^n  y_i - \hat{y}_i $  | <code>loss = "mae"</code>                      |
| Binary<br>Cross Entropy      | $L = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i))$ | <code>loss = "binary_crossentropy"</code>      |
| Multi-Class<br>Cross Entropy | $L = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K y_{ik} \log(\hat{p}_{ik})$                | <code>loss = "categorical_crossentropy"</code> |

- ▶ Cross entropy is the appropriate loss function for the classification.

# Keras Basics (5/9)

| Gradient descent algorithms:

| Algorithm                   | Explanation   | Keras  |
|-----------------------------|---|--|
| Stochastic Gradient Descent | Gradient descent algorithm with batch learning.         | <code>SGD(lr=0.1)</code>                               |
| Momentum                    | Takes into account the gradient from the previous step. | <code>SGD(momentum=0.9)</code>                         |
| Nesterov Momentum           | Minimizes unnecessary movements.                        | <code>SGD(lr=0.01, momentum=0.9, nesterov=True)</code> |
| Adagrad                     | Adaptative step size.                                   | <code>Adagrad(lr=0.01)</code>                          |
| RMSProp                     | Improves upon the Adagrad.                              | <code>RMSProp(lr=0.01)</code>                          |
| Adam                        | Combines the best properties of Adagrad and RMSProp.    | <code>Adam(lr=0.01)</code>                             |

# Keras Basics (6/9)

Activation functions:

| Name    | Formula  | Keras                             |
|---------|--|-----------------------------------|
| Linear  | $linear(x) = x$                                      | <code>activation="linear"</code>  |
| Sigmoid | $\sigma(x) = \frac{1}{1 + e^{-x}}$                   | <code>activation="sigmoid"</code> |
| Tanh    | $tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$        | <code>activation="tanh"</code>    |
| ReLU    | $ReLu(x) = \max(0, x)$                               | <code>activation="relu"</code>    |
| Softmax | $\sigma(x)_j = \frac{e^{x_j}}{\sum_{k=1}^K e^{x_k}}$ | <code>activation="softmax"</code> |

# Keras Basics (7/9)

## Layers:

| Keras             | Explanation  |
|-------------------|--|
| Dense()           | Dense layer.   |
| Conv1D()          | 1D convolution layer.  |
| Conv2D()          | 2D convolution layer.  |
| MaxPooling1D()    | 1D max pooling layer.  |
| MaxPooling2D()    | 2D max pooling layer.  |
| Dropout()         | Applies dropout to the input.                                |
| Flatten()         | Flattens the input.  |
| Embedding()       | Turns positive integers (label encoding) into dense vectors. |
| SimpleRNN()       | RNN where the output is fed back to input.                   |
| LSTM()            | LSTM layer.  |
| TimeDistributed() | A wrapper that applies a layer to every time slice.          |
| Input()           | The first layer for a functional API model.                  |

# Keras Basics (8/9)

| Two ways of building models with Keras:

1). Sequential API:

- Layers are added to a "Sequential" object.
- Easy but less flexible.

2). Functional API:

- The first layer must be defined as "Input" where the data shape is specified.
- Each layer defined as a function that takes in as argument the previous layer.
- A "Model" object takes in as arguments the first and the last layers.
- More flexible.

# Keras Basics (9/9)

## I Datasets provided by Keras:

- Useful for learning/practicing purpose.
- Provided by the "keras.datasets" module.

| Name           | Description   |
|----------------|---|
| cifar10        | 32 × 32 color images labeled over 10 categories. 50000 training and 10000 testing images.   |
| cifar100       | 32 × 32 color images labeled over 100 categories. 50000 training and 10000 testing images.  |
| imdb           | 25000 movie reviews from IMDB with binary labeling (positive or negative).                  |
| reuters        | 11228 newswires from Reuters, labeled over 46 topics.                                       |
| mnist          | 28 × 28 grayscale images of handwritten digits. 60000 training and 10000 testing images.    |
| fashion_mnist  | 28 × 28 grayscale images of 10 fashion categories. 60000 training and 10000 testing images. |
| boston_housing | 13 attributes of houses at different locations around Boston suburbs in the late 1970s.     |

# Deep Learning - Part II

## UNIT 3. Deep Learning with Keras.

3.1. Keras Basics.

### **3.2. AI with Keras.**

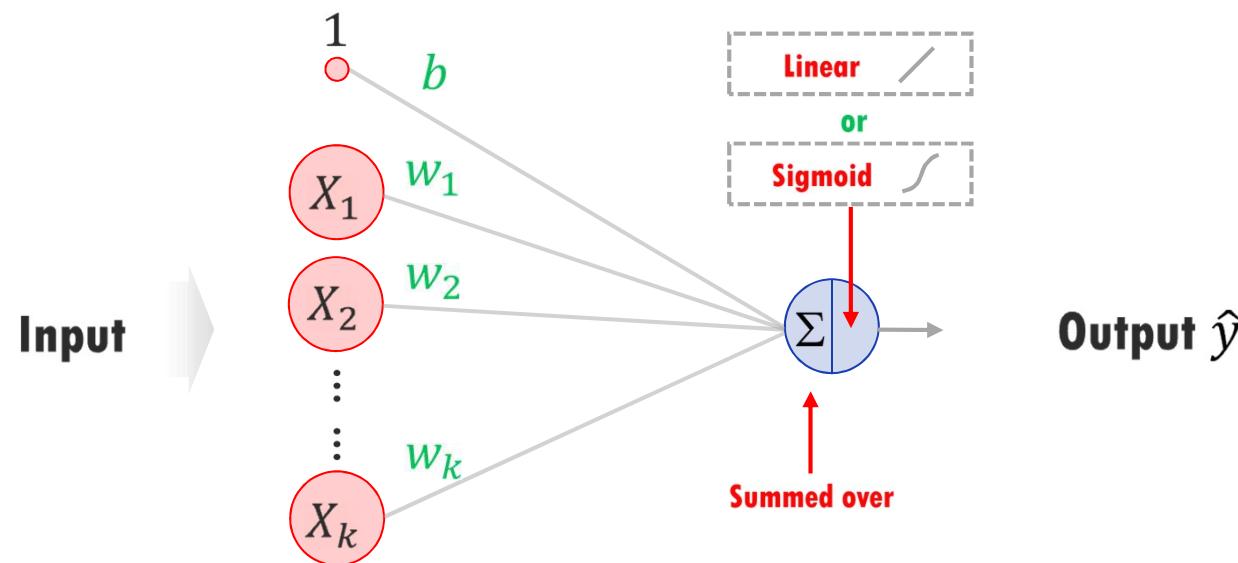
3.3. Natural Language Processing with Keras.

**UNIT 3.**

**3.2. AI with Keras.**

# AI with Keras (1/8)

| Linear/Logistic regression with Keras:



- ▶ The major difference between linear and logistic regression lies in the choice of activation function.

**UNIT 3.**

**3.2. AI with Keras.**

# AI with Keras (2/8)

## | Linear/Logistic regression with Keras:

- ▶ Sequential API example:

```
# Import the necessary classes.  
from keras.models import Sequential  
from keras.layers import Dense
```

```
# Build the model by adding a layer.  
my_model = Sequential()  
my_model.add(Dense(input_dim = n_vars, units = 1, activation="linear")) # activation="linear" for linear regression.
```

```
# Define the optimization method and the loss function. Then, compile.  
my_optimizer=Adam(lr=learn_rate)  
my_model.compile(loss = "mse", optimizer = my_optimizer, metrics=["mse"])
```

```
# Train.  
my_model.fit(X_train, y_train, epochs=n_epochs, batch_size = n_batch_size, validation_split = 0.2, verbose = 0)
```

## UNIT 3.

### 3.2. AI with Keras.

# AI with Keras (3/8)

## | Linear/Logistic regression with Keras:

- ▶ Functional API example:

```
# Import the necessary classes.  
from keras.models import Model  
from keras.layers import Input, Dense
```

```
# Build the model by adding a layer.  
my_input = Input(shape=(n_vars,)) # The first layer with the shape specified.  
my_output = Dense(units=1,activation="linear")(my_input) # Takes in the previous layer as function argument.  
my_model = Model(inputs=my_input, outputs=my_output) # The first and the last layers entered.
```

- ▶ Optimization, compilation, training, etc. steps are the same as before.

# Coding Exercise #0801

**Follow practice steps on ‘ex\_0801.ipynb’ file.**

# Coding Exercise #0802

**Follow practice steps on ‘ex\_0802.ipynb’ file.**

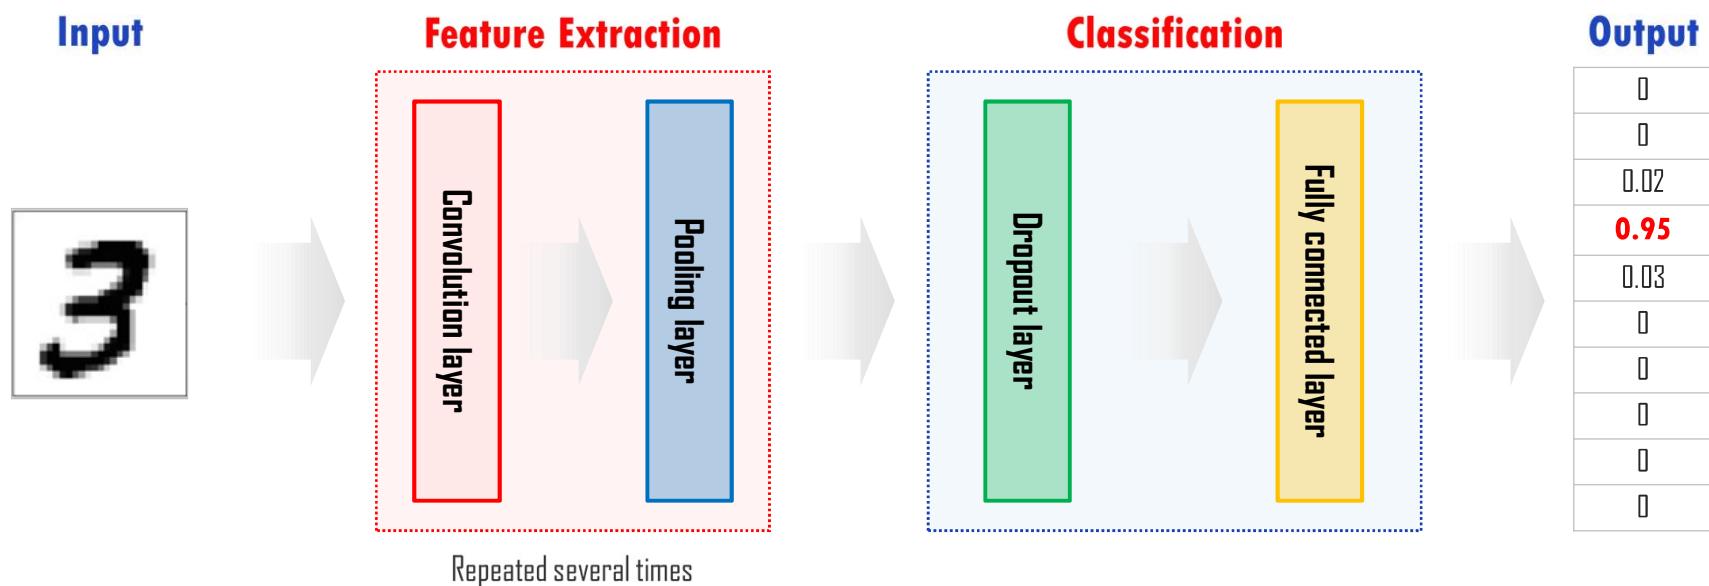
**UNIT 3.**

**3.2. AI with Keras.**

# AI with Keras (4/8)

| Diagram of a Convolutional Neural Network (CNN):

- The convolution and pooling layers are repeated several times.
- There is a fully connected layer just before the output.



**UNIT 3.**

**3.2. AI with Keras.**

# AI with Keras (5/8)

| CNN with Keras: a code example.

```
# Import the necessary classes.  
from keras.models import Sequential  
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D  
# We will use the Sequential API.
```

```
# Build the CNN model by adding the layers.  
hold_prob = 0.5  
my_model = Sequential()  
my_model.add(Conv2D(input_shape=(28,28,1),filters=32,kernel_size=(6,6), activation="relu"))  
my_model.add(MaxPooling2D(pool_size=2,padding='same'))  
my_model.add(Conv2D(filters=64,kernel_size=(6,6), activation="relu"))  
my_model.add(MaxPooling2D(pool_size=2,padding='same'))  
my_model.add(Flatten())  
my_model.add(Dense(units = 1024, activation="relu"))  
my_model.add(Dropout(rate=hold_prob))  
my_model.add(Dense(units = 10, activation="softmax"))
```

# Coding Exercise #0803

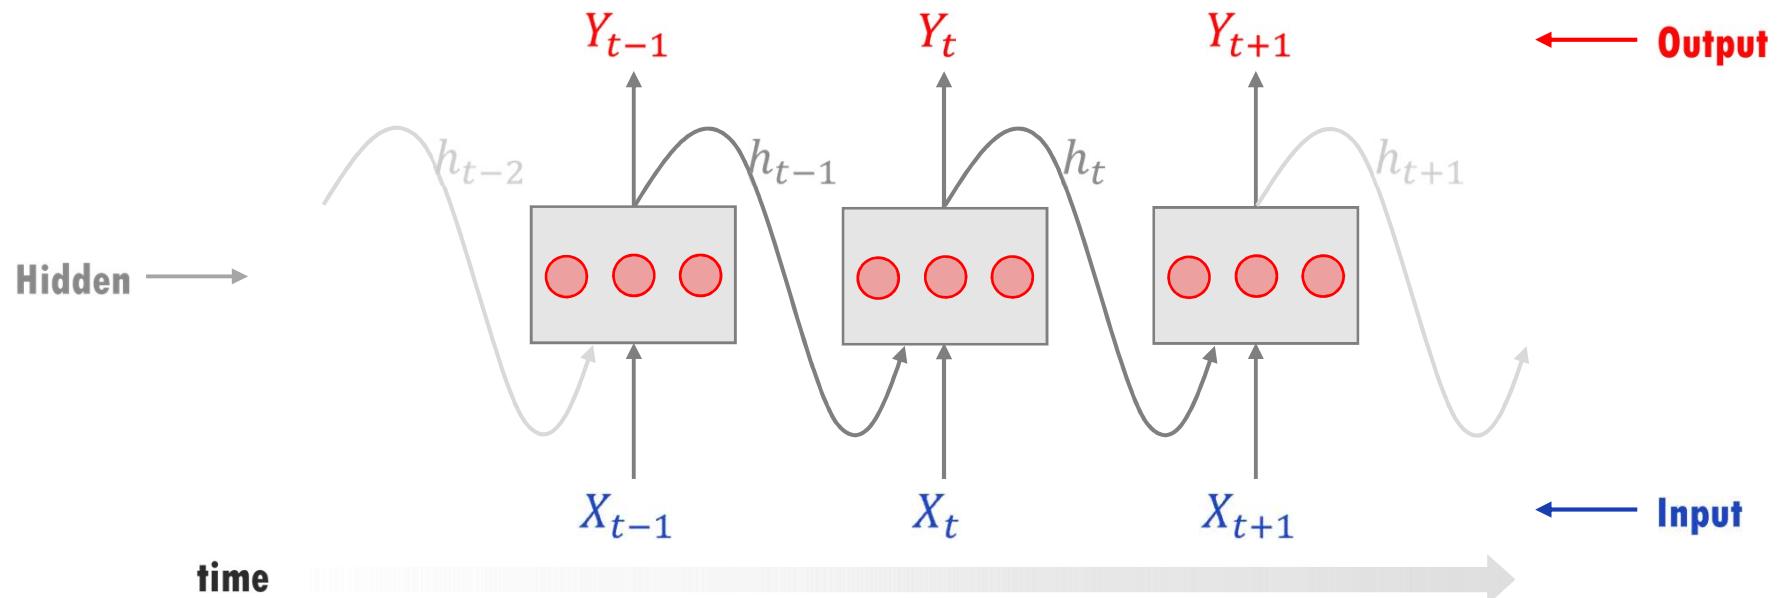
**Follow practice steps on ‘ex\_0803.ipynb’ file.**

# Coding Exercise #0804

**Follow practice steps on ‘ex\_0804.ipynb’ file.**

# AI with Keras (6/8)

| Recurrent Neural Network (RNN):



- This is a "Sequence in and Sequence out" model.

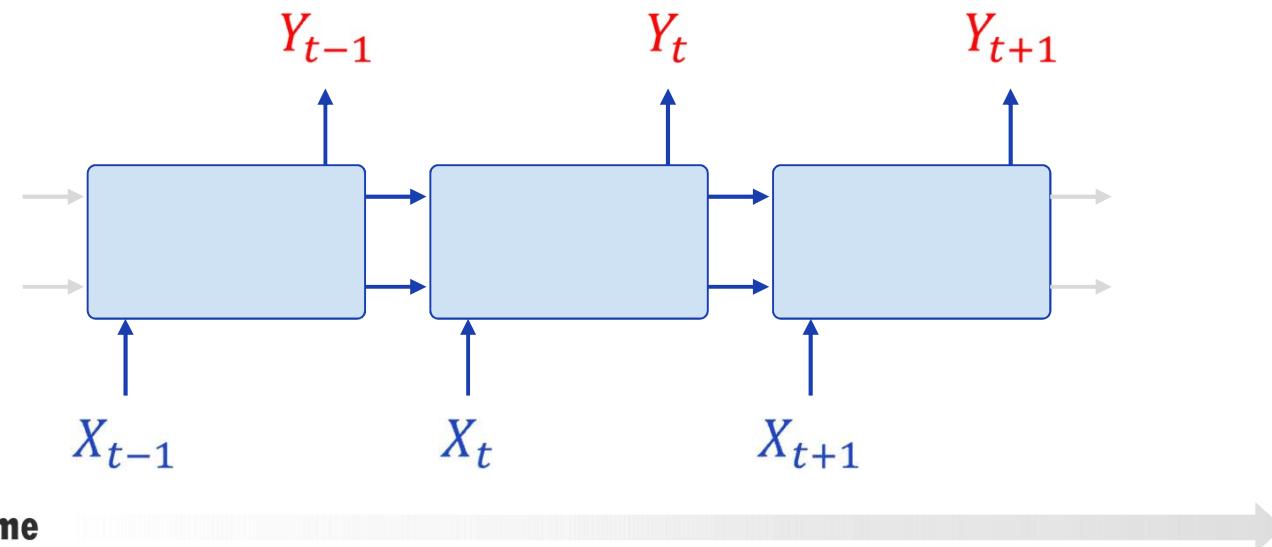
**UNIT 3.**

**3.2. AI with Keras.**

# AI with Keras (7/8)

## | Long Short Term Memory (LSTM) network:

- LSTM cells can form a sequence just like the regular RNN cells.
- The following is a "Sequence in and Sequence out" model.



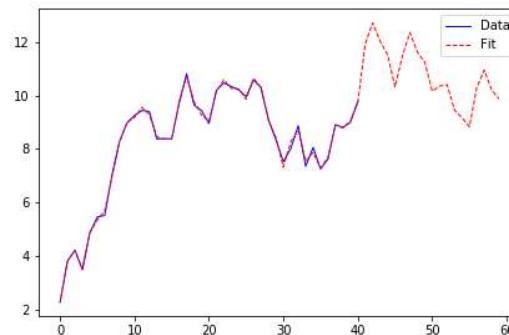
## UNIT 3.

### 3.2. AI with Keras.

# AI with Keras (8/8)

I RNN/LSTM with Keras: a code example.

```
# Import the necessary classes.  
from keras.models import Sequential  
from keras.layers import SimpleRNN, LSTM, Dense, TimeDistributed  
  
# "sequence in and sequence out" model : at SimpleRNN() return_sequences=True.  
my_model = Sequential()  
my_model.add(SimpleRNN(units=n_neurons, return_sequences=True, input_shape=(None, n_input))) # RNN.  
my_model.add(TimeDistributed(Dense(units=n_output, activation="linear"))) # Wrapped.
```



# Coding Exercise #0805

**Follow practice steps on ‘ex\_0805.ipynb’ file.**

# Deep Learning - Part II

## UNIT 3. Deep Learning with Keras.

3.1. Keras Basics.

3.2. AI with Keras.

**3.3. Natural Language Processing with Keras.**

# Natural Language Processing with Keras (1/16)

## One-hot-encoding vs word embedding:

**Ex**). Given a sentence "I eat an apple every morning", let's suppose that the words are indexed as:

|         |   |   |
|---------|---|---|
| I       | : | 3 |
| Eat     | : | 0 |
| An      | : | 2 |
| Apple   | : | 1 |
| Every   | : | 4 |
| Morning | : | 5 |

The words would have the following one-hot-encoding representations:

|         |   |               |
|---------|---|---------------|
| I       | : | [0 0 0 1 0 0] |
| Eat     | : | [1 0 0 0 0 0] |
| An      | : | [0 0 1 0 0 0] |
| Apple   | : | [0 1 0 0 0 0] |
| Every   | : | [0 0 0 0 1 0] |
| Morning | : | [0 0 0 0 0 1] |

# Natural Language Processing with Keras (2/16)

## One-hot-encoding vs word embedding:

- We notice obvious problems with the one-hot-encoding representation.
- To improve, the “word embedding” is introduced which is a distributed representation method.

| One-Hot-Encoding   | Embedding  |
|--|--|
| The dimension of the vector space is large.<br>The dimension is as large as the vocabulary size. | The dimension of the vector space is limited.                        |
| Vectors are sparse; they are mostly filled with 0s that carry no information.                    | Vectors are dense.<br>Every vector element carries some information. |
| No semantic relationship among the vectors.<br>The vectors are orthogonal to each other.         | Semantic relationship among the vectors.                             |

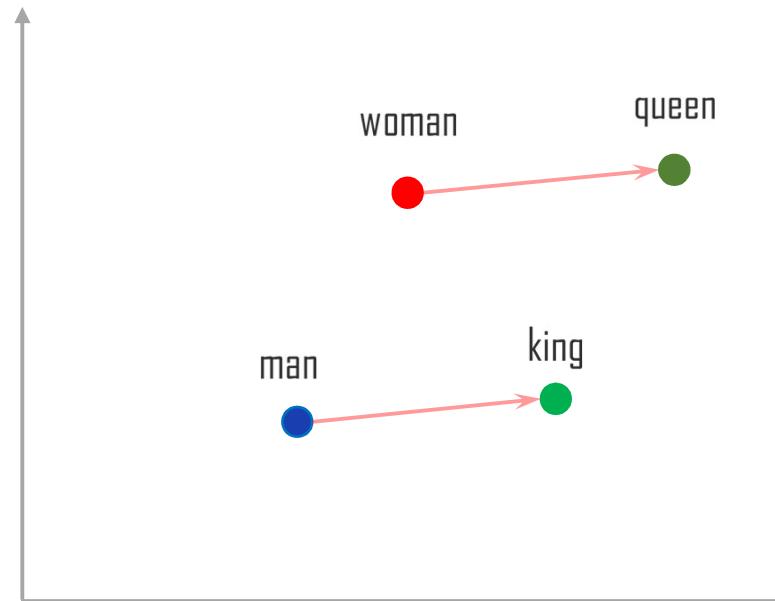
- There are also “paragraph embedding” and “document embedding” representations.
- We will call “dense vector” or “embedding vector” interchangeably.

**UNIT 3.**

**3.3. Natural Language Processing with Keras.**

# Natural Language Processing with Keras (3/16)

| Word embedding (Word2Vec):



- Among the dense vectors, relationships such as following are established:

$$\text{queen} - \text{woman} = \text{king} - \text{man}$$

# Natural Language Processing with Keras (4/16)

## | Word embedding (Word2Vec):

- Use CBOW (Continuous Bag of Words) and/or Skip-Gram to build the embedding vectors.
  - 1). Build a predictive model based on the Softmax regression (multi-class logistic regression).
  - 2). We assume one-hot-encoded input and output vectors.
  - 3). Extract the embedding vectors from the trained weight matrices.

## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (5/16)

## | Word embedding (Word2Vec):

- ▶ CBOW: Using the context words, predict the (missing) center word.

| Training Sentence            | Center Word | Context Words  |
|------------------------------|-------------|----------------|
| I eat an apple every morning | I           | eat            |
| I eat an apple every morning | eat         | I, an          |
| I eat an apple every morning | an          | eat, apple     |
| I eat an apple every morning | apple       | an, every      |
| I eat an apple every morning | every       | apple, morning |
| I eat an apple every morning | morning     | every          |

- ▶ We assumed a "sliding window" over the training sentence.

# Natural Language Processing with Keras (6/16)

## | Word embedding (Word2Vec):

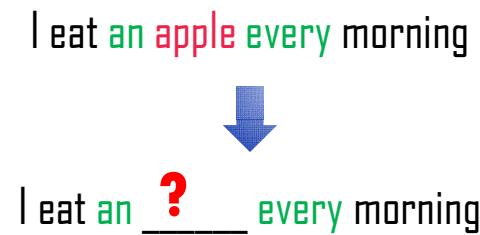
- CBOW: Using the context words, predict the (missing) center word.

**Ex).** Let's suppose that the vector dimension of the one-hot-encoded words = 6.

Let's also suppose that we would like to find dense vectors of dimension = 3.

We will consider two context words (one from the left and another from the right).

So, we have a situation as the following:



## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (7/16)

## | Word embedding (Word2Vec):

- CBOW: Using the context words, predict the (missing) center word.

**Ex).** (continues from the previous page)

Then, we build a Softmax regression model.

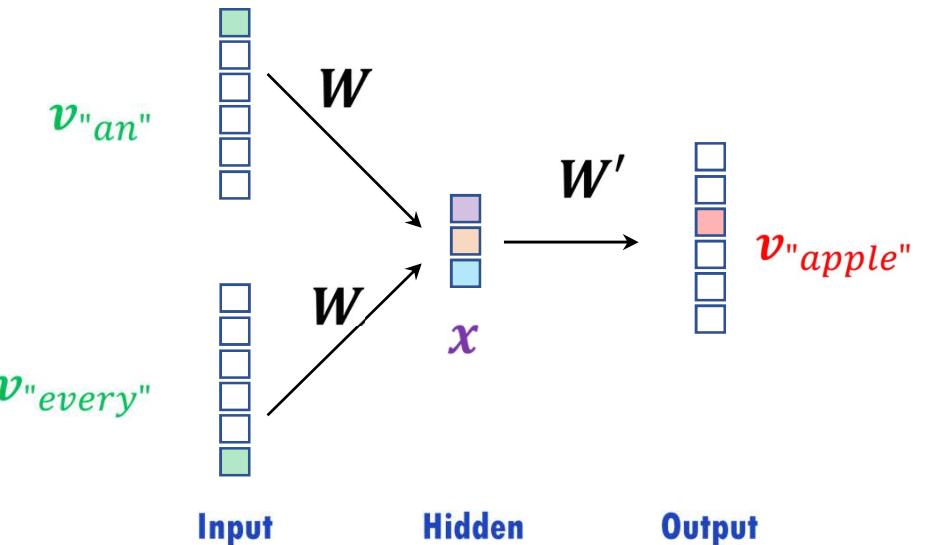
For the vector inputs of "an" and "every",

we would like to train the weights  $\mathbf{W}$

and  $\mathbf{W}'$  such that the predicted output is

the vector "apple".

One-hot-encoded words :  $\mathbf{v}^{\text{"an"}} = [1 \ 0 \ 0 \ 0 \ 0 \ 0]$  ,  $\mathbf{v}^{\text{"every"}} = [0 \ 0 \ 0 \ 0 \ 0 \ 1]$  ,  $\mathbf{v}^{\text{"apple"}} = [0 \ 0 \ 1 \ 0 \ 0 \ 0]$



# Natural Language Processing with Keras (8/16)

## | Word embedding (Word2Vec):

- CBOW: Using the context words, predict the (missing) center word.

**Ex).** (continues from the previous page)

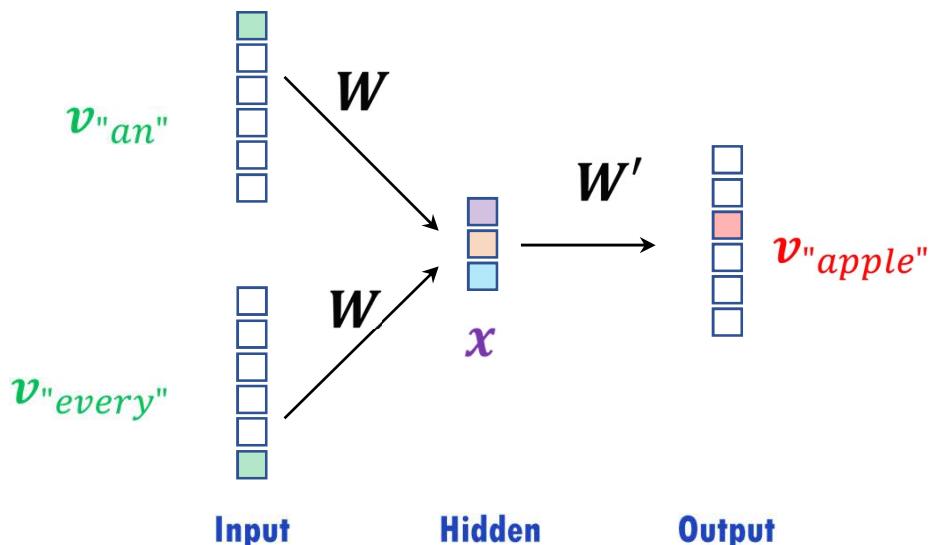
We have the following sizes:

Size of the matrix  $W=3 \times 6$

Size of the matrix  $W'=6 \times 3$

Dimension of the vector  $x=3$ .

Dimension of the input and output = 6.



## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (9/16)

## | Word embedding (Word2Vec):

- CBOW: Using the context words, predict the (missing) center word.

**Ex.** (continues from the previous page)

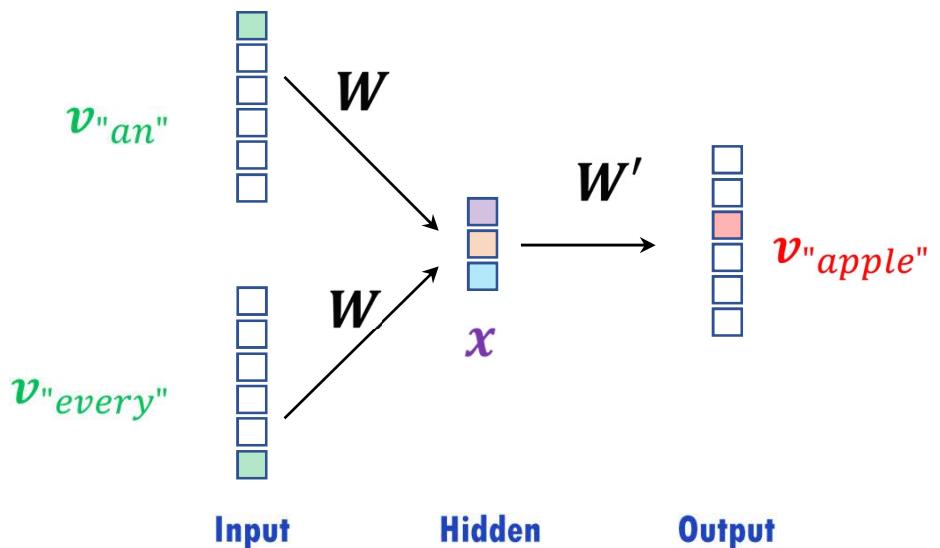
We propagate forward from the input layer

to the hidden layer (a single node):

$$x^{an} = W \cdot v^{an}$$

$$x^{every} = W \cdot v^{every}$$

$$x = \frac{x^{an} + x^{every}}{2} \quad \Leftarrow \text{Average for the hidden node.}$$



## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (10/16)

## | Word embedding (Word2Vec):

- CBOW: Using the context words, predict the (missing) center word.

**Ex.** (continues from the previous page)

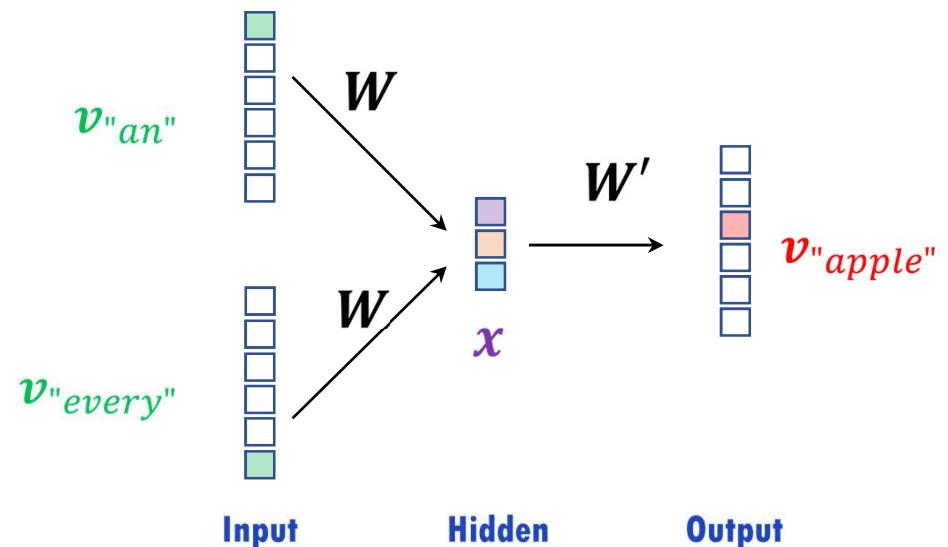
We propagate forward to the output layer:

$$\hat{v} = W' \cdot x$$

We should train the weights  $W$  and  $W'$

$\text{argmax}(\hat{v})$  and  $\text{argmax}(v^{\text{"apple"}}$ )

is minimized.



# Natural Language Processing with Keras (11/16)

## | Word embedding (Word2Vec):

- CBOW: Using the context words, predict the (missing) center word.

**Ex).** (continues from the previous page)

Now, let's interpret the result.

a). When we propagate from the input layer to the hidden layer (by matrix multiplication),

the one-hot-encoded input vectors  $v_{\text{"an"}}$  and  $v_{\text{"every"}}$  are picking the columns 0 and 5

of  $W$  and projecting them to the hidden layer with the target dimension = 3.

b). So, the dense vectors for "an", "every" are the **columns** 0 and 5 of the **trained**  $W$ .

c). Analogously, we can extract dense vectors from the **rows** of the **trained**  $W'$ .

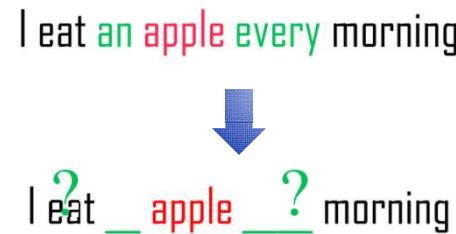
**UNIT 3.**

### 3.3. Natural Language Processing with Keras.

## Natural Language Processing with Keras (12/16)

### I Word embedding (Word2Vec):

- ▶ Skip-Gram: Using a center word, predict the (missing) context words.



I eat **an** **apple** **every** morning

↓

I **eat** \_ **apple** \_ **?** morning

- Similar to the CBOW, here also we train a Softmax regression to predict the missing words.
- We extract the dense vectors (embedding vectors) from the trained weight matrices.

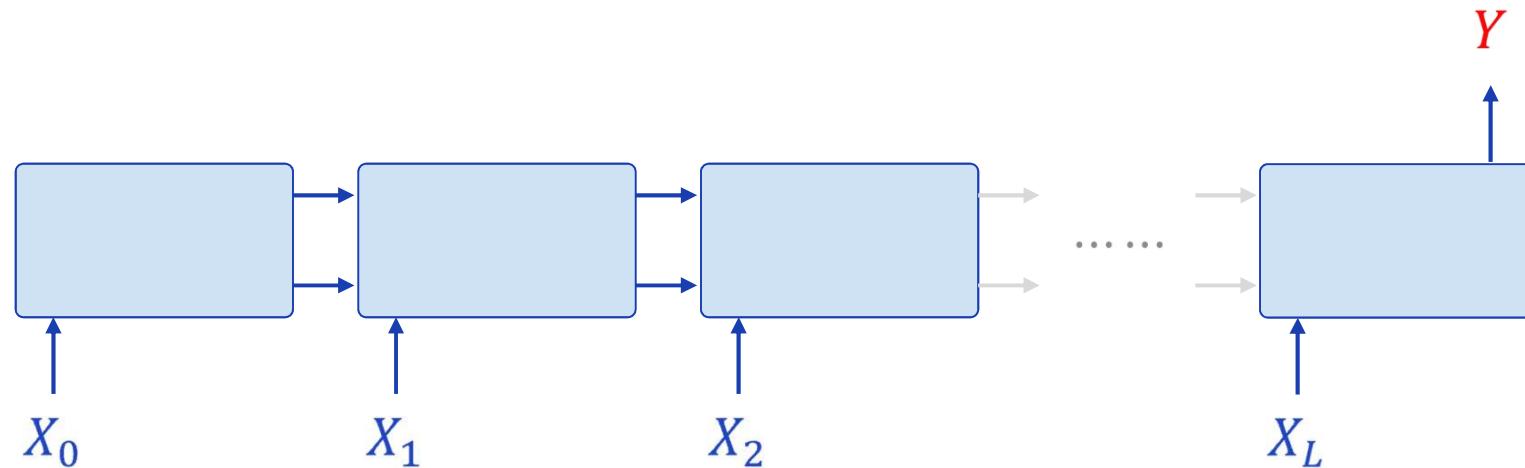
**UNIT 3.**

**3.3. Natural Language Processing with Keras.**

# Natural Language Processing with Keras (13/16)

I LSTM network for document classification:

- "Sequence in and Vector out" model.
- Embedding representation of the words.



# Natural Language Processing with Keras (14/16)

I LSTM network for document classification: a code example.

```
# Import the necessary classes.  
from keras.models import Sequential  
from keras.layers import Dense, LSTM, Embedding  
# We will use the Sequential API.
```

```
# Build a model by adding the layers.  
my_model = Sequential()  
my_model.add(Embedding(n_words,n_input))  
my_model.add(LSTM(units=n_neurons, return_sequences=False, input_shape=(None, n_input), activation='tanh'))  
my_model.add(Dense(1, activation='sigmoid'))
```

- ▶ In `LSTM()`, we should set `return_sequences=False` for a "Sequence in and Vector out" model.

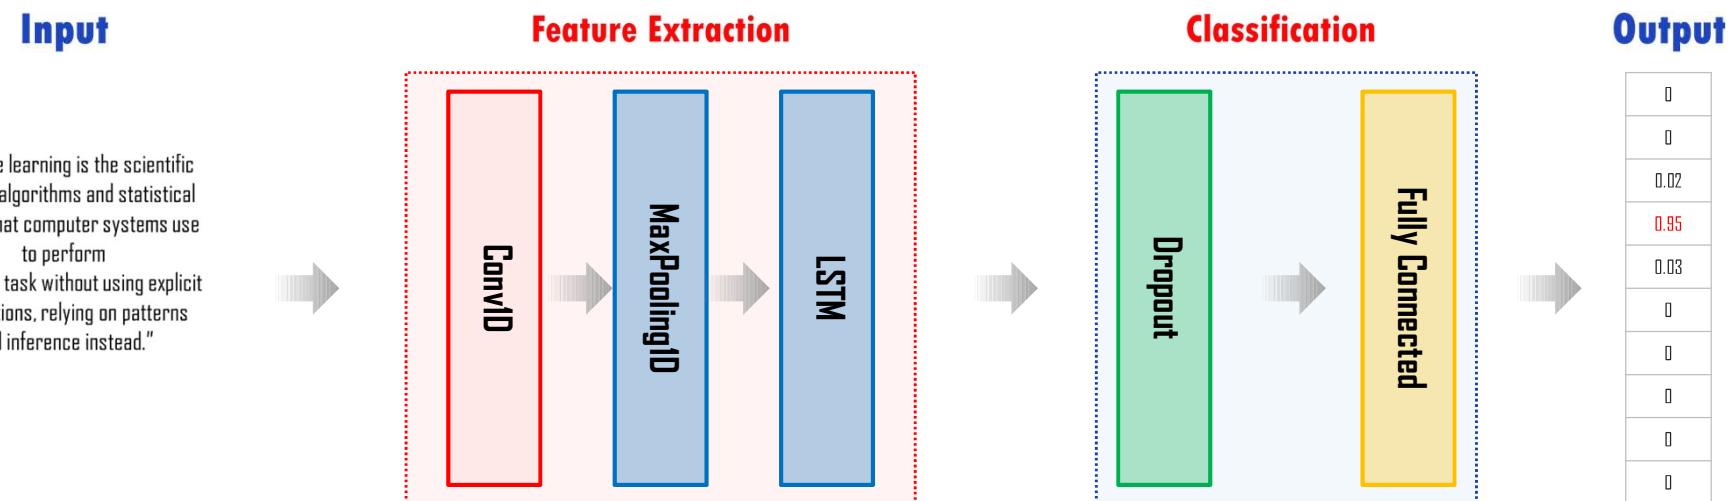
## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (15/16)

I Deep learning model for document classification:

- 1D convolution + 1D max pooling + LSTM for the feature extraction.
- Localized sequence patterns picked up by the 1D convolution.



## UNIT 3.

### 3.3. Natural Language Processing with Keras.

# Natural Language Processing with Keras (16/16)

I Deep learning model for document classification: a code example.

```
# Import the necessary classes.  
from keras.models import Sequential  
from keras.layers import Dense, LSTM, Embedding, Conv1D, MaxPool1D, Dropout  
# We will use the Sequential API.
```

```
# Build a model by adding the layers.  
my_model = Sequential()  
my_model.add(Embedding(n_words, n_input)) # Embedding layer.  
my_model.add(Conv1D(filters=n_filters, kernel_size = k_size, strides=stride_size,padding='valid',activation='relu'))  
my_model.add(MaxPool1D(pool_size = 2))  
my_model.add(LSTM(units=n_neurons, return_sequences=False, input_shape=(None, n_input), activation='tanh'))  
my_model.add(Dropout(rate=hold_prob))  
my_model.add(Dense(l, activation='sigmoid'))
```

# Coding Exercise #0806

**Follow practice steps on ‘ex\_0806.ipynb’ file.**

# Coding Exercise #0807

**Follow practice steps on ‘ex\_0807.ipynb’ file.**

# Coding Exercise #0808

**Follow practice steps on ‘ex\_0808.ipynb’ file.**

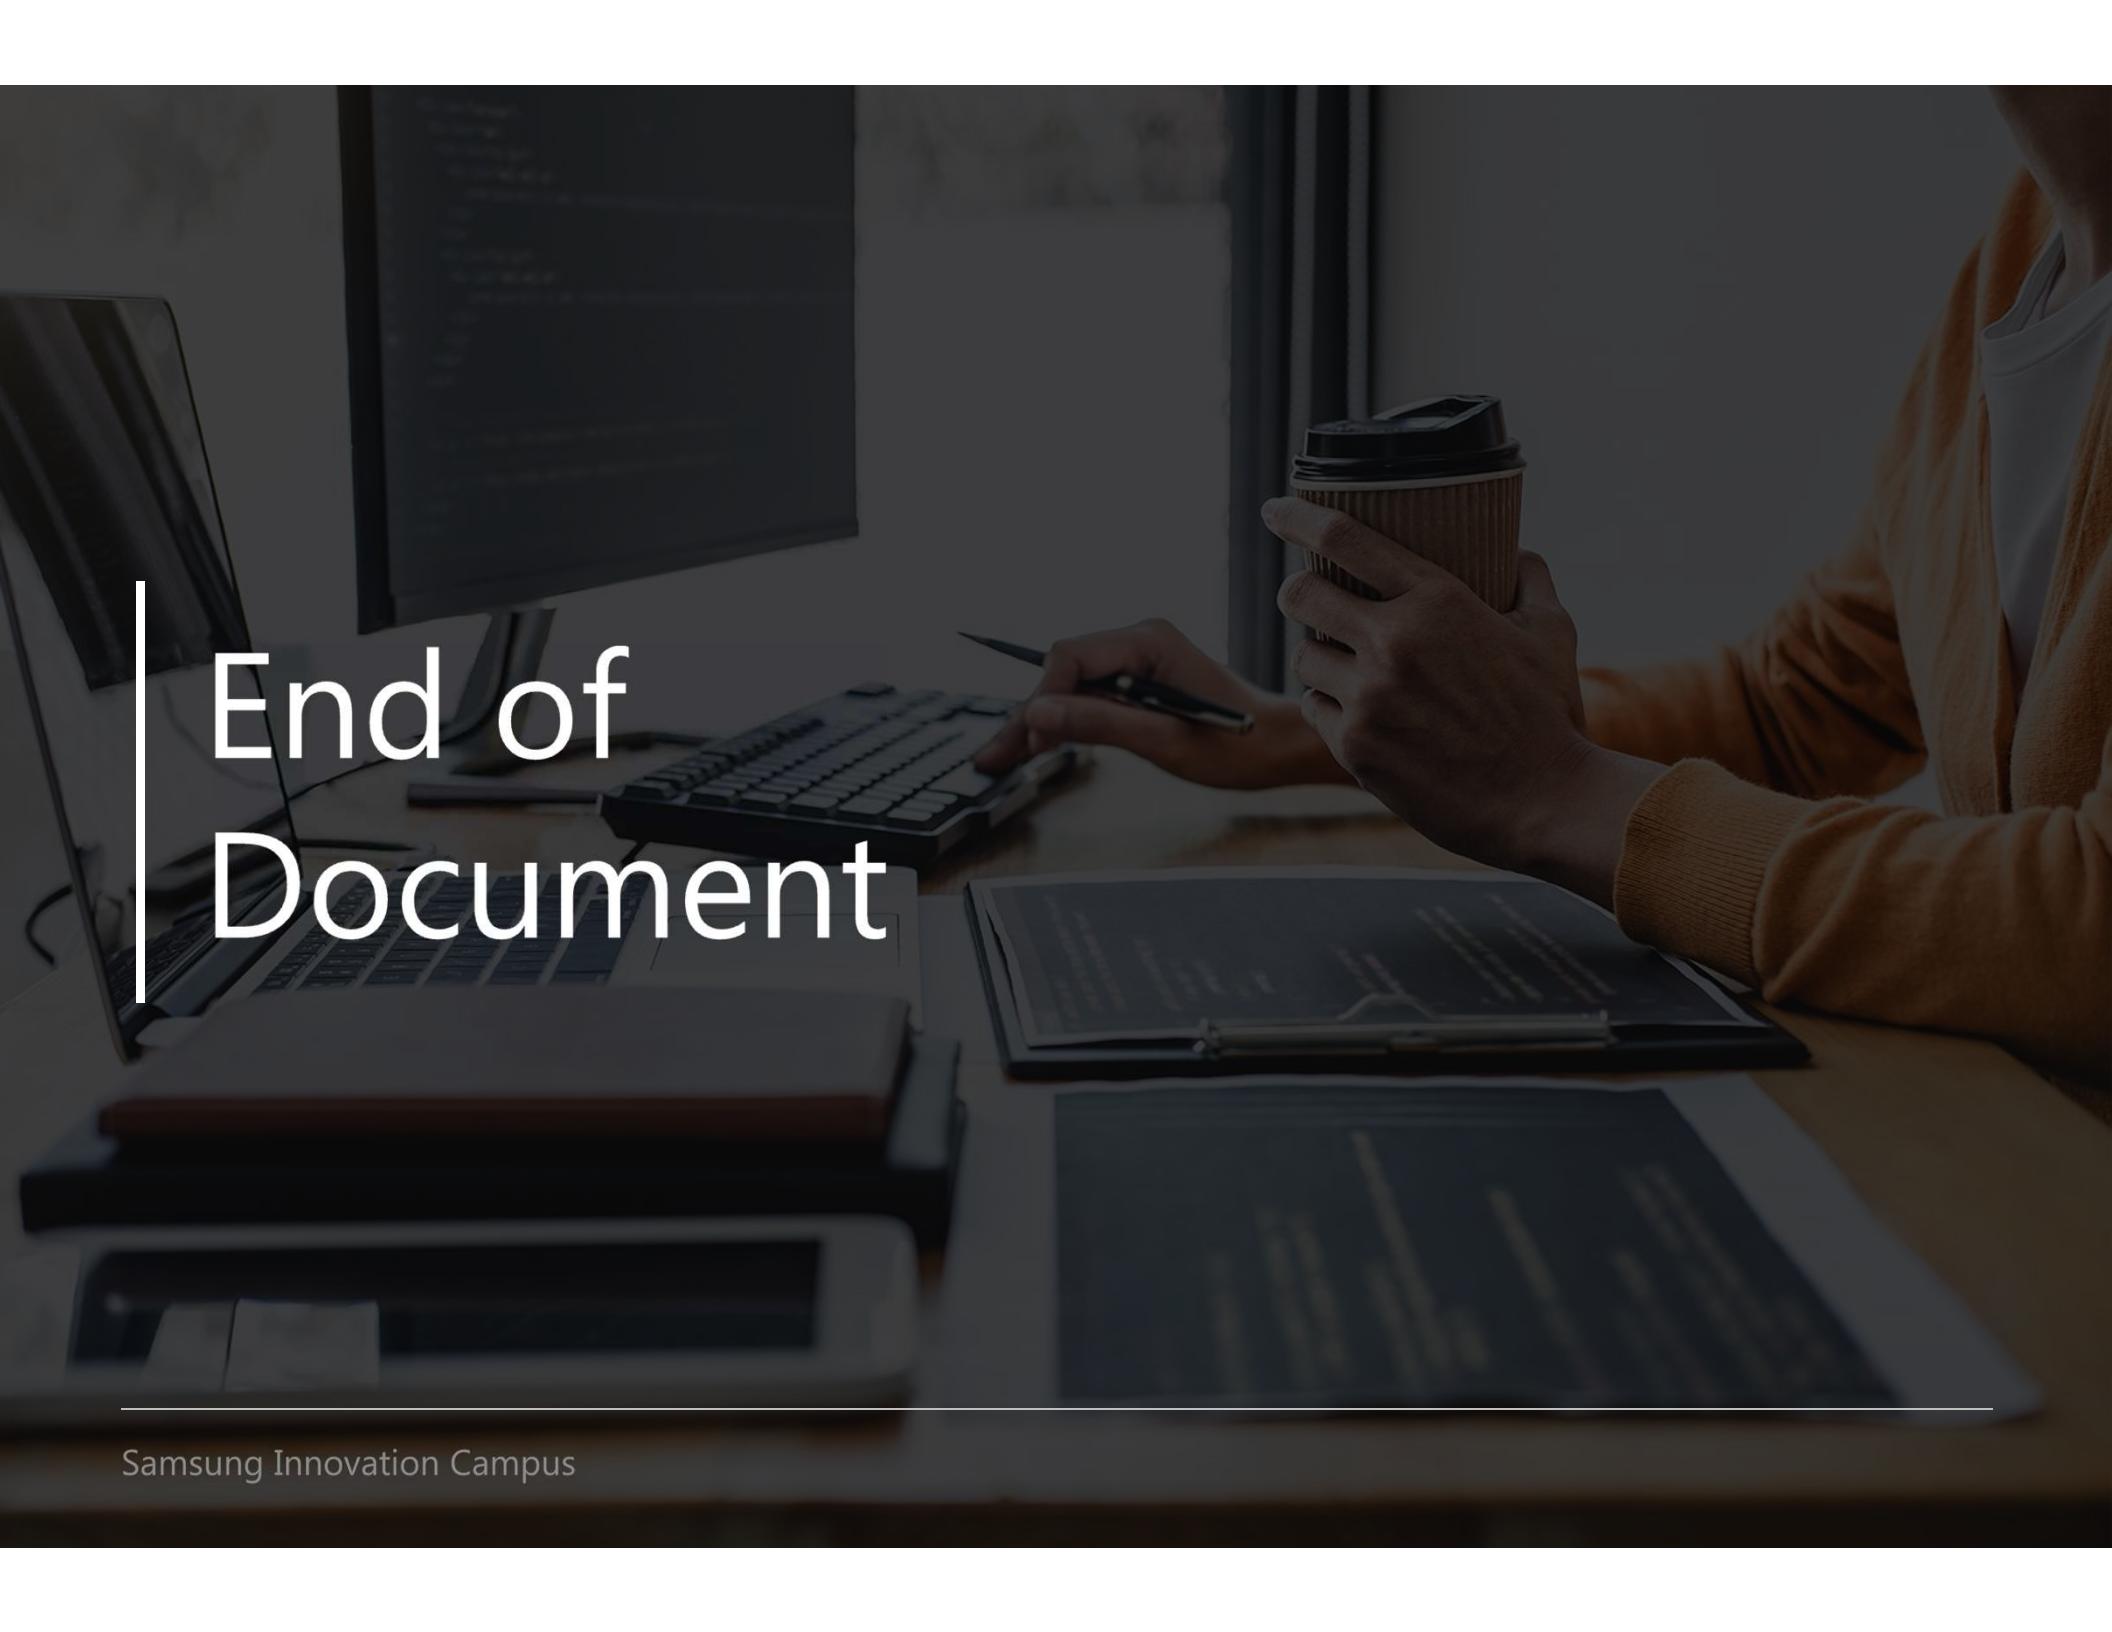
# Coding Exercise #0809

**Follow practice steps on ‘ex\_0809.ipynb’ file.**

End of chapter Quiz

# Quiz #0801

Duration : 2 Hours

A photograph of a person working at a desk. They are wearing a light-colored long-sleeved shirt and are holding a brown paper coffee cup with a black lid in their right hand. Their left hand is on a keyboard. In front of them is a laptop with its screen open. To the left of the laptop, there is a stack of papers or books. The background shows two computer monitors on stands. The overall scene suggests a professional or academic environment.

# End of Document

Together for Tomorrow!  
**Enabling People**

Education for Future Generations

©2019 SAMSUNG. All rights reserved.

Samsung Electronics Corporate Citizenship Office holds the copyright of book.

This book is a literary property protected by copyright law so reprint and reproduction without permission are prohibited.

To use this book other than the curriculum of Samsung innovation Campus or to use the entire or part of this book, you must receive written consent from copyright holder.